DATA WORKS SUMMIT

#DSW2018

10 Steps to Make Your Cluster Hum

Alejandro Fernandez

San Jose
June 17-21, 2018

Speaker

Alejandro Fernandez

airbnb

Infrastructure Engineer

’17 San Jose

’16 Melbourne

’16 San Jose

PMC

Apache Ambari

Contributed

HIVE

Superset

Airflow

Poll

Managing your Hadoop cluster?

Easy

Medium

Difficult

Who is this for?

DevOps & Infra Teams
Technical audience

Dev: Add features

QE: Test for security & scale

Op: Deployment + business continuity

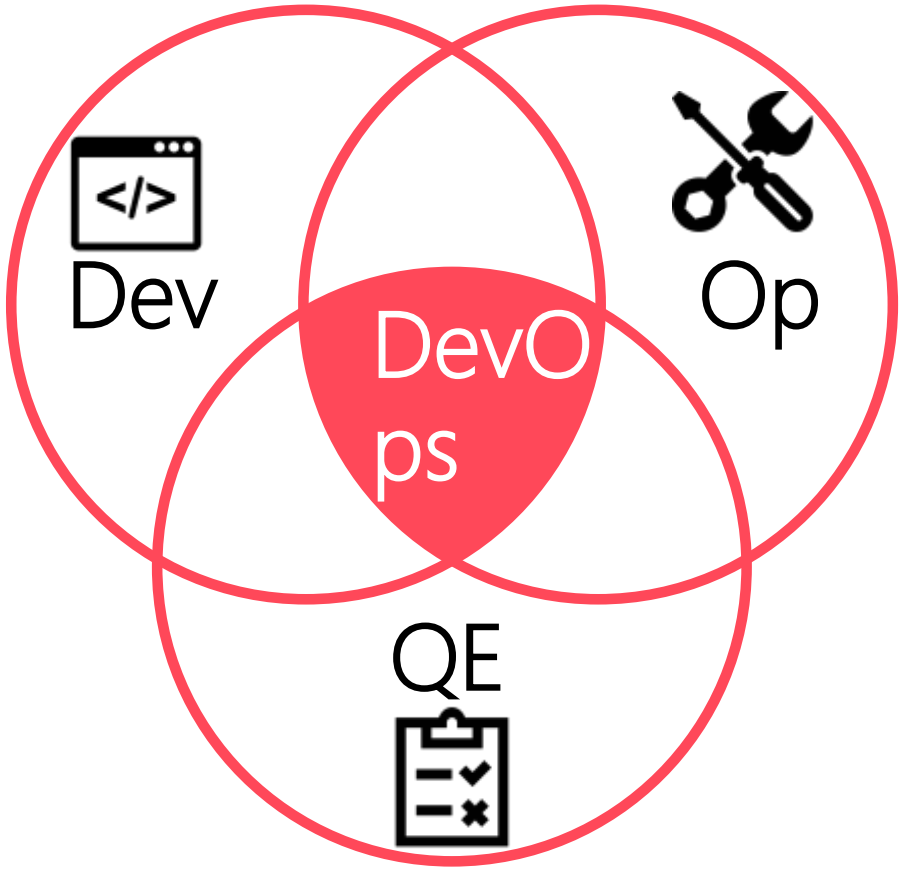

Steps To a Smoother Cluster

1. Automate adding hosts
2. Host Name discovery
3. Artifact management
4. Deployments and upgrades
5. Log searching with ELK

Steps To a Smoother Cluster

6. Monitor Hadoop metrics with Ambari Metrics System
7. Data-driven (metrics + alerts)
8. Cron-jobs, ETL with Apache Airflow (incubating)
9. Visualize trends with Apache Superset (incubating)
10. Automated maintenance

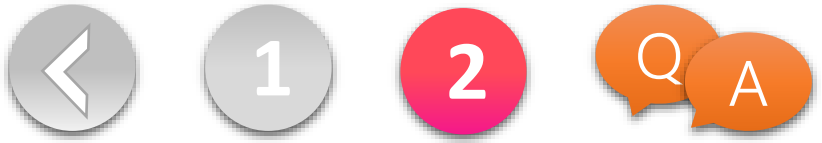

1. Deploy New Hosts

- Using Chef
  - Cookbooks (collection of templates, attributes, and recipes)
  - Recipes (patterns + blocks)
  - Roles (NM, DN, NM, etc.)
- Repeatable
- Idempotent + fixable

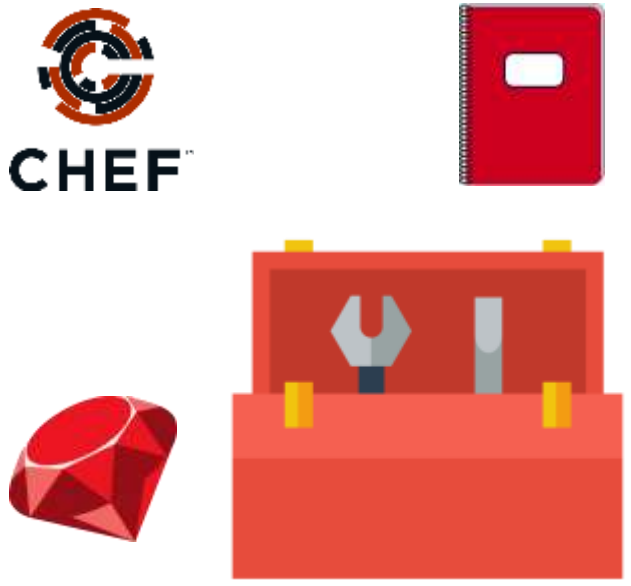

Automatically Add Hosts to Cluster

- Add hosts to cluster
- Assign roles
- Rack awareness
- Host discovery

Apache Ambari

curl http://

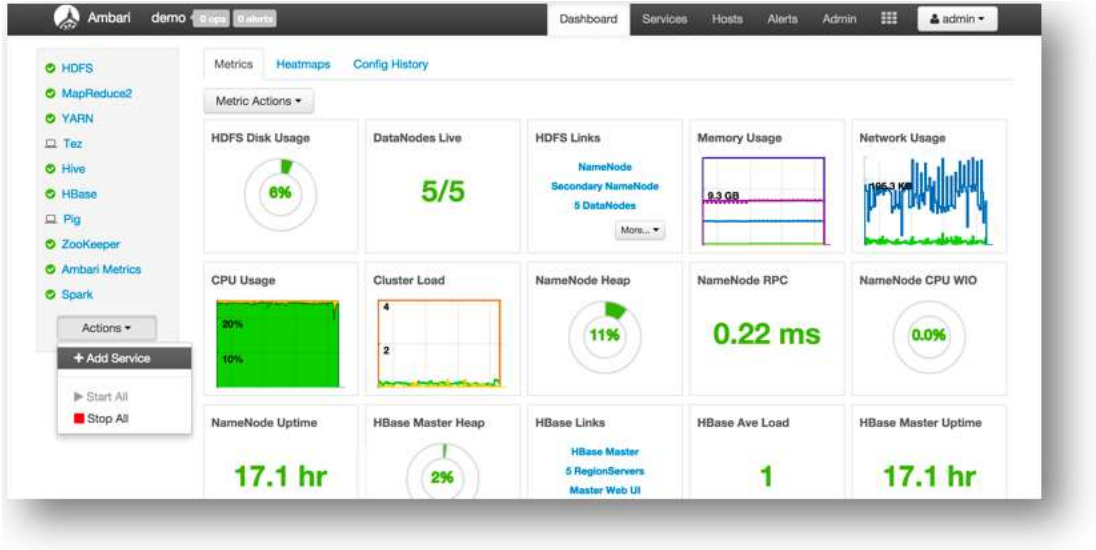

Cloudera Manager

API

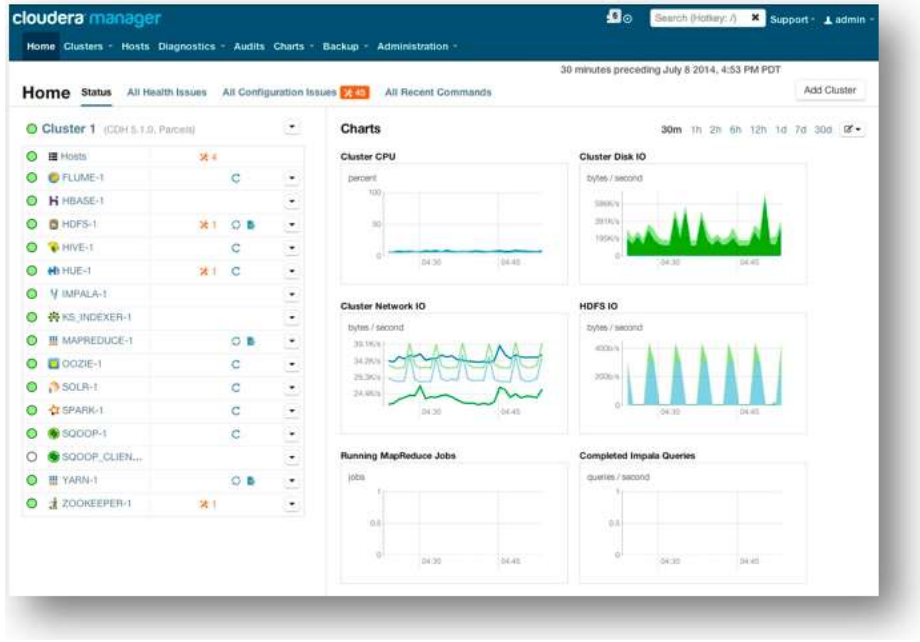

2. Host Name Discovery

Avoid hardcoding hostnames in configs

File	Config
yarn-site.xml	hadoop.registry.zk.quorum
hbase-site.xml	hbase.zookeeper.quorum
hive-site.xml	hive.metastore.uris hive.zookeeper.quorum
storm-site.xml	nimbus.seeds
etc	etc

- Replace bad host
- ~~Change hard-coded configs~~
- ~~Restart services dependent on those configs~~

Name Discovery Solutions

Method	Pros	Cons
Apache Ambari Stack Advisor	Works out of the box. {{ var }} is defined in params.py	Write your own for custom services
Chef Attributes	Can handle multiple environments with same recipe	Configs are written to disk, but not using CM or Ambari
Smartstack	Local load-balanced servers behind the same name	

3. Trust The Artifacts

Artifact: build, jars, scripts, resources, build log, and test results

- Reproducible: Generated by build system
- Trackable: Embed version, Git SHA that generated it
- Trustworthy: Hash and Signed

Simple

Enterprise

Deployments

1. Artifact server

2. Always take backups

3. Deploy parcels or bits
   - Orchestration: rolling restart if possible
   - Verification: analyze alerts on hosts with new changes
   - Downgrade: ready to rollback

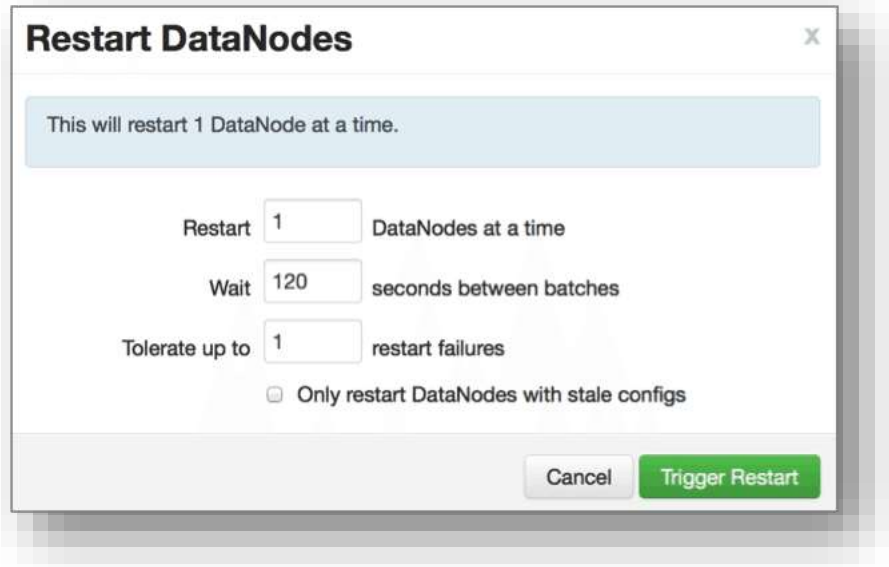

Automate Calls to Upgrade

- Prechecks:
  - All services up
  - Ran Service Checks after config changes
  - All hosts heartbeating
  - No hosts in maintenance mode

  …
- Start: Options for failure handling
- Status

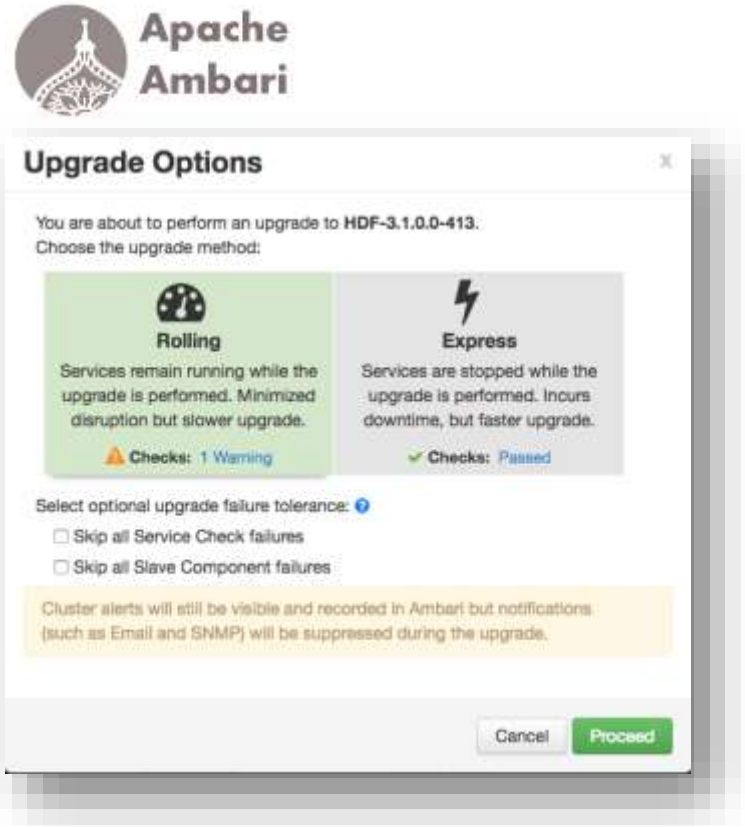

5. ELK Stack for Searching Logs

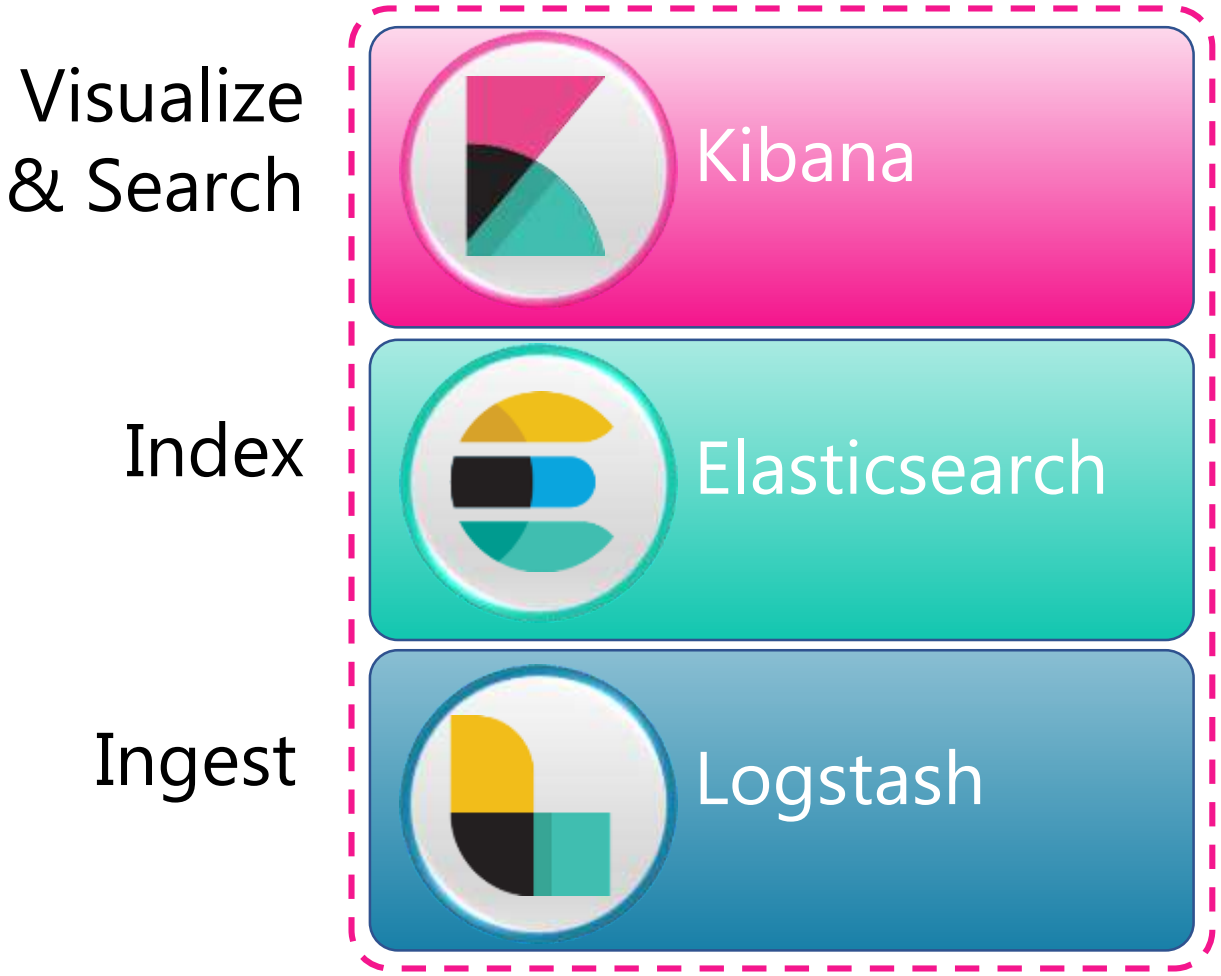

Out-of-the-Box Log Search

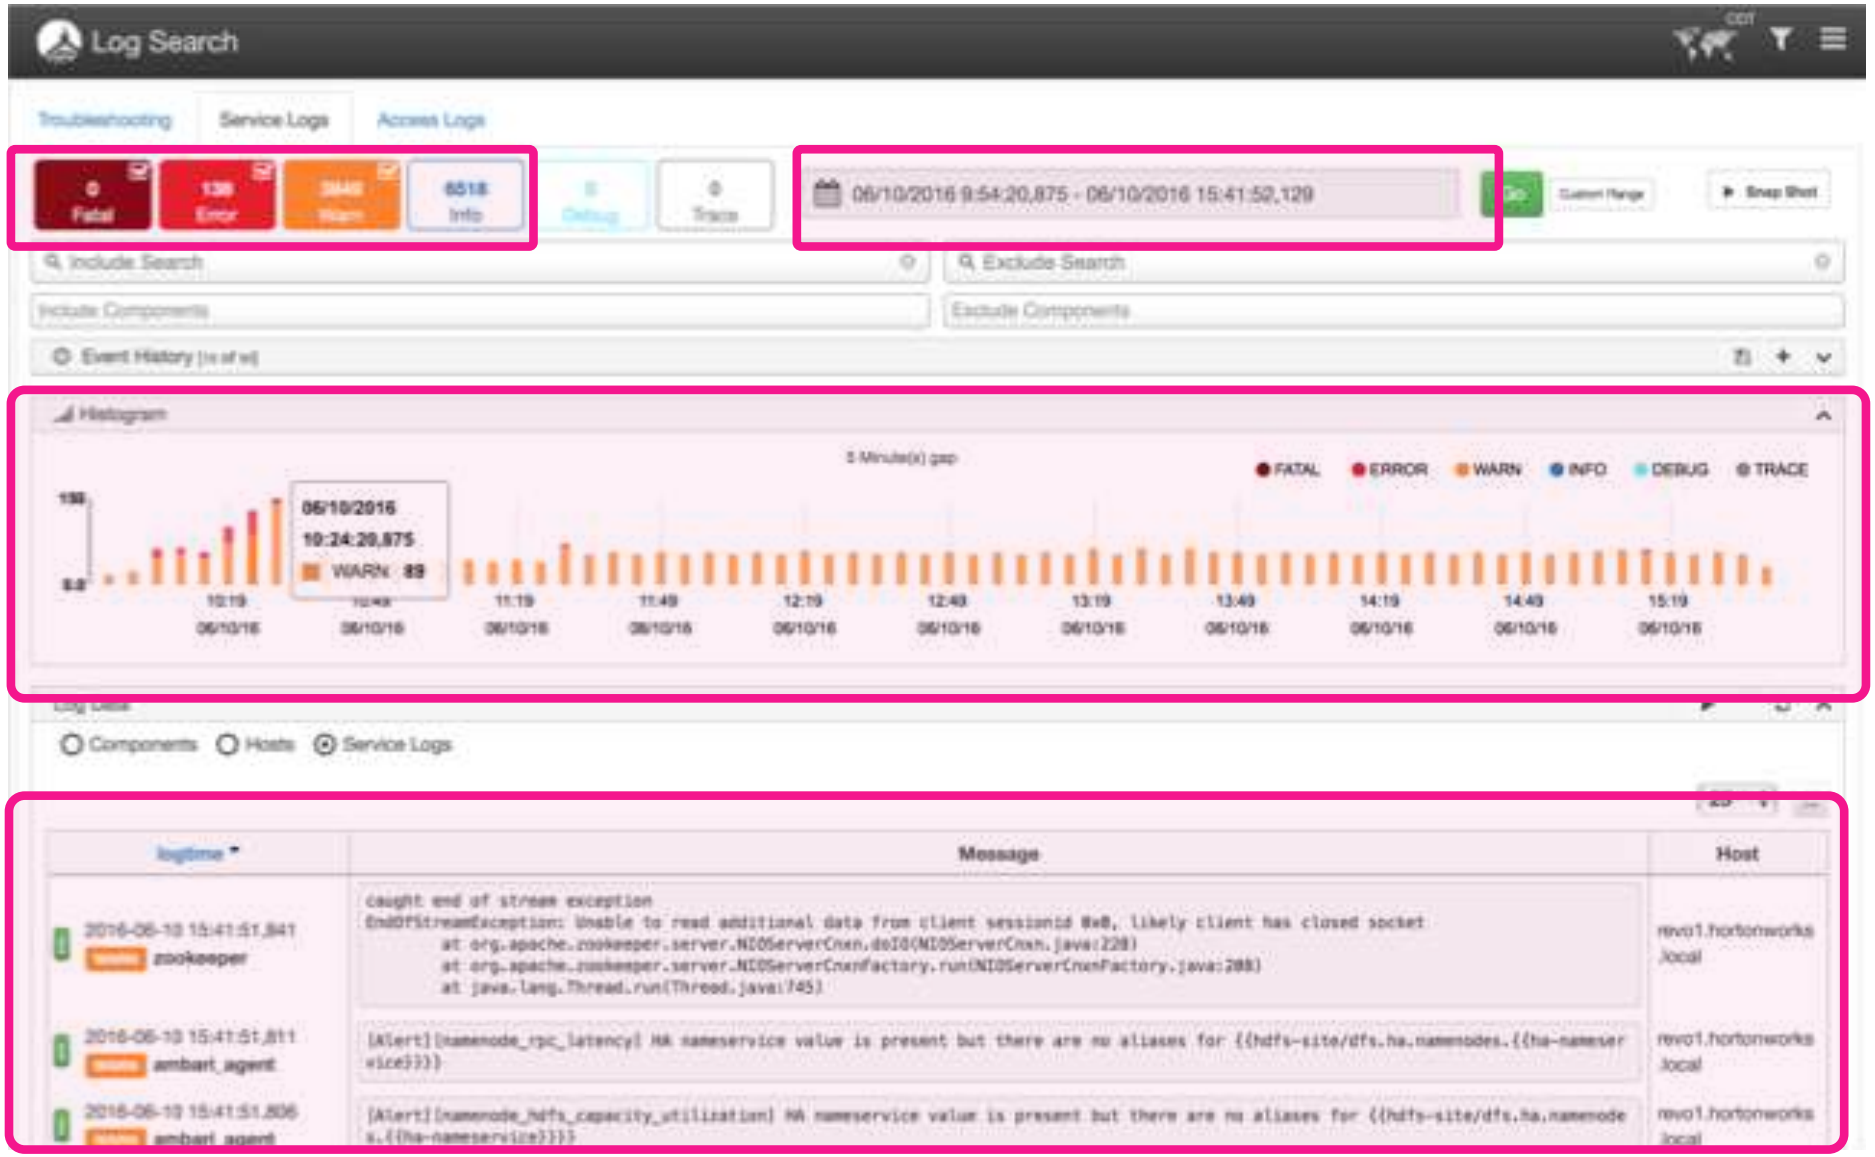

Advanced Log Search

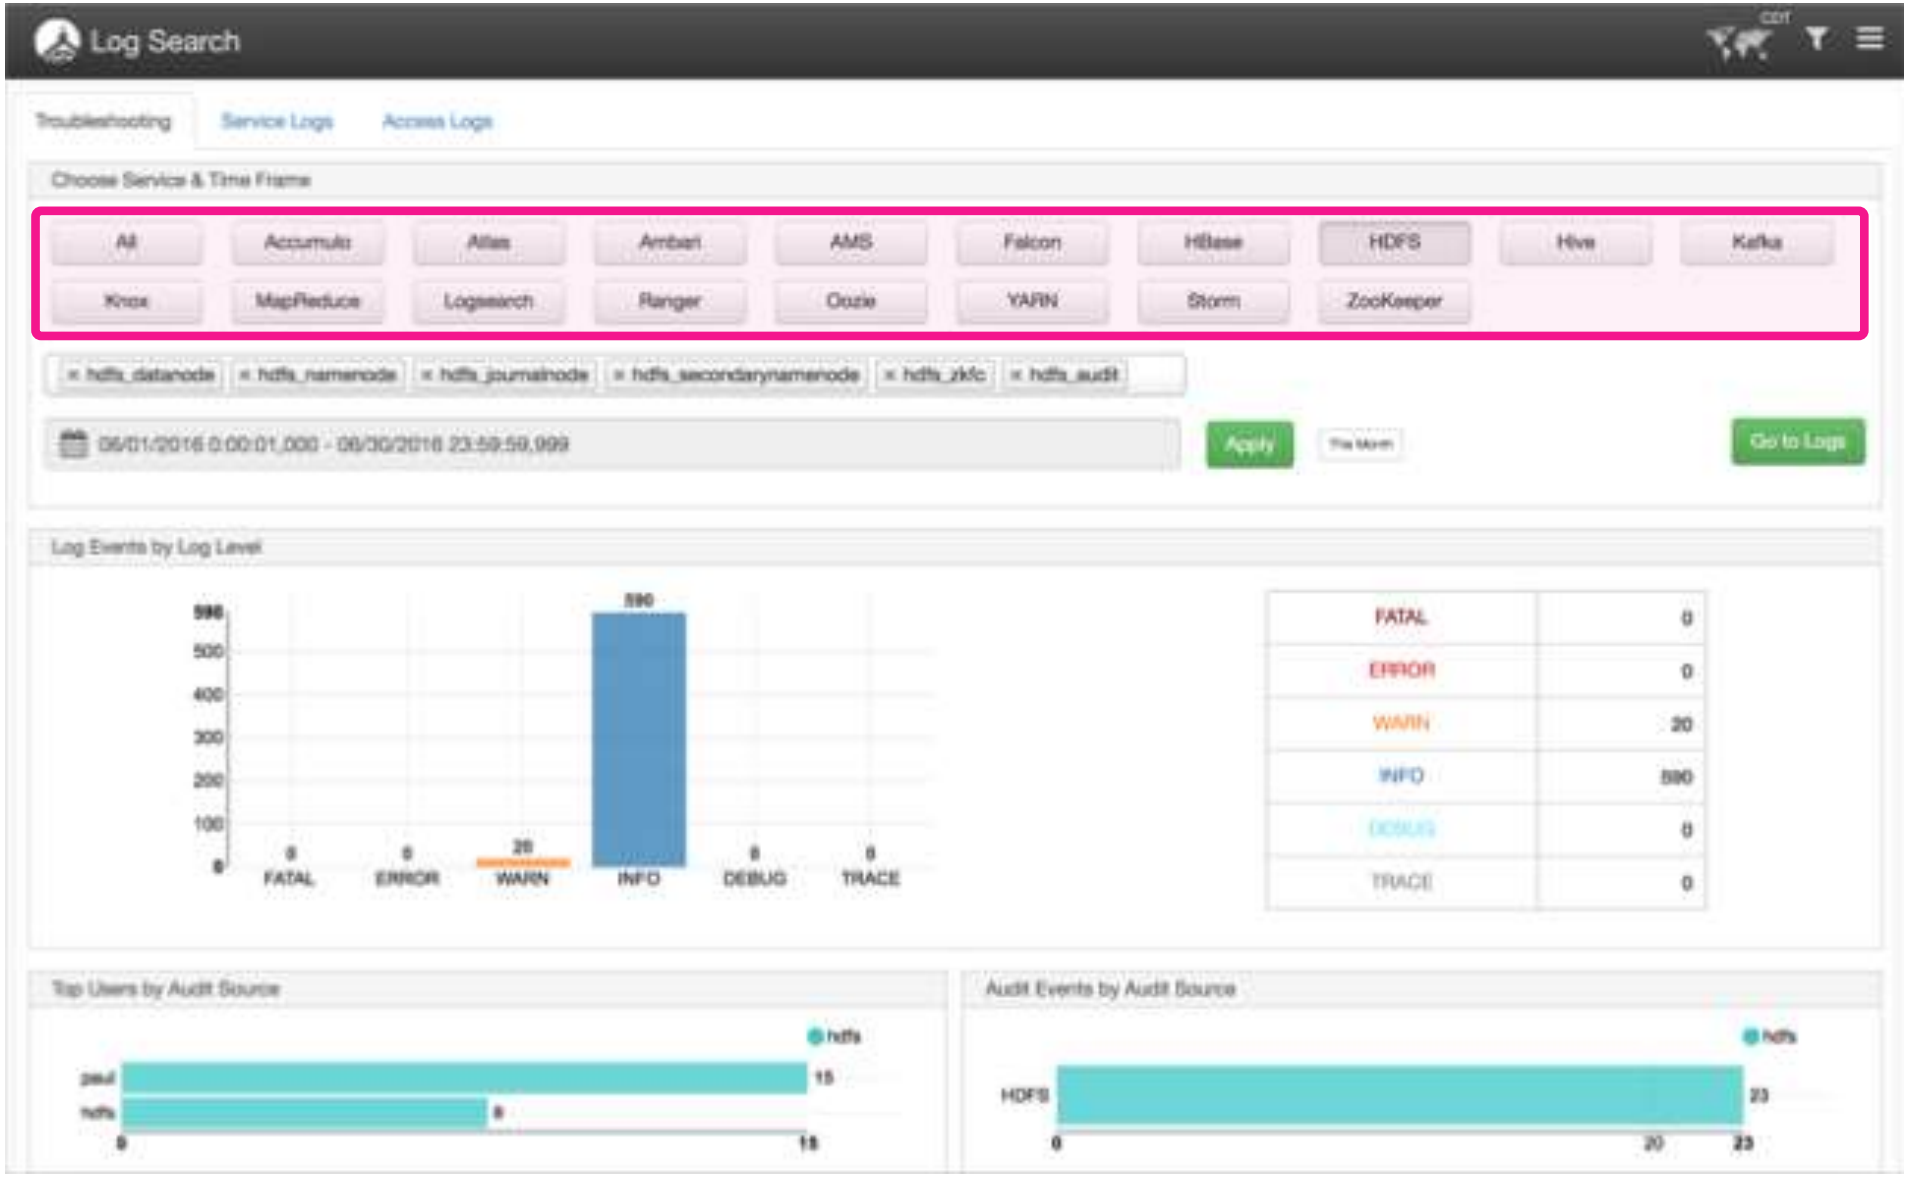

Advanced Log Search

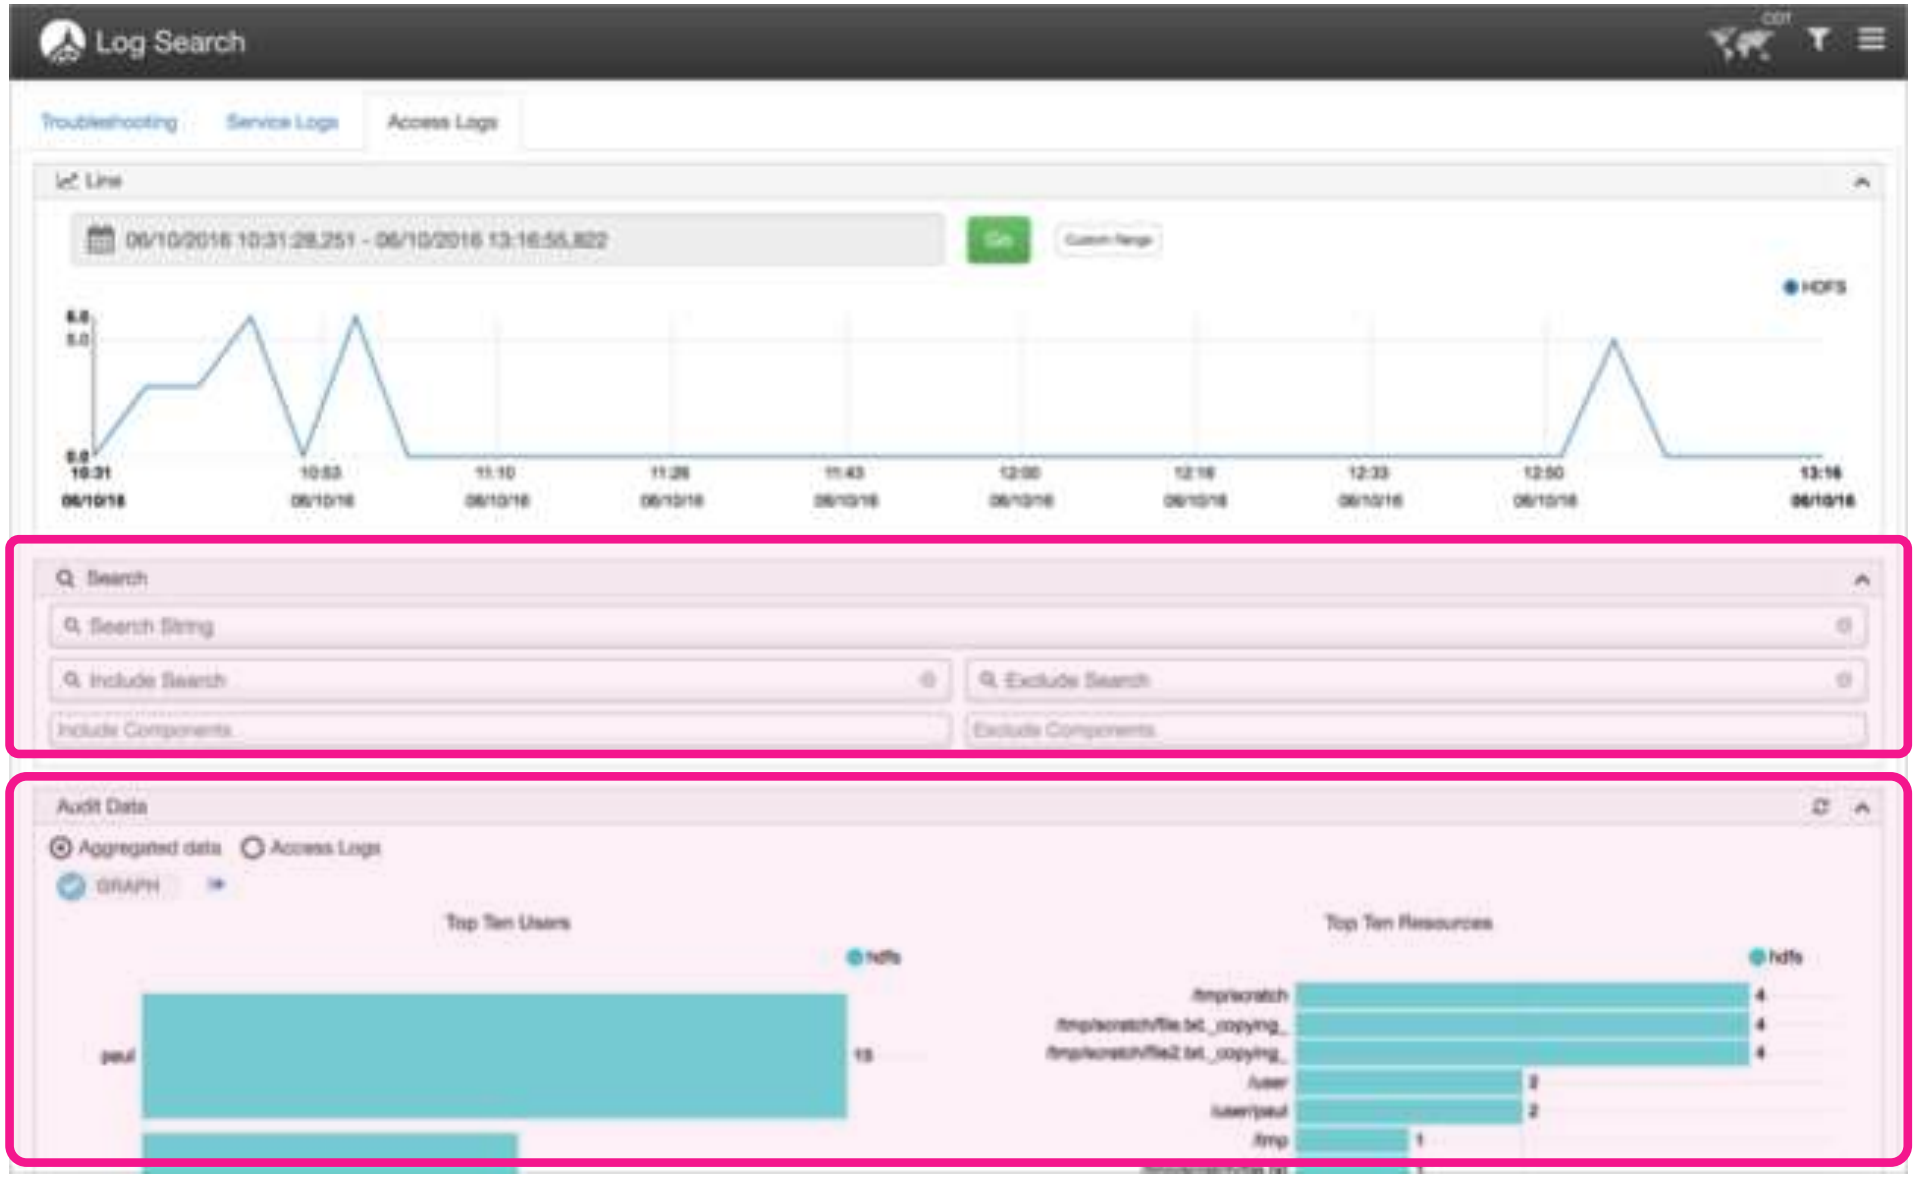

6. Collecting Metrics

- Apache Zookeeper:
  - Jmxterm
  - Inspect the beans: num connections, latency, timeouts, etc.

- Apache HDFS and YARN:
  - WebHDFS stats on a directory
  - ResourceManager web UI job counters
  - JMX

and more…

```
curl http://$HOST:$PORT/webhdfs/v1
```

6. Collecting with AMS

Ambari Metrics System for Hadoop metrics:

- Support Hadoop stack using Hadoop syncs interface
- Anomaly Detection
- Query with Phoenix
- High Availability
- Scale horizontally

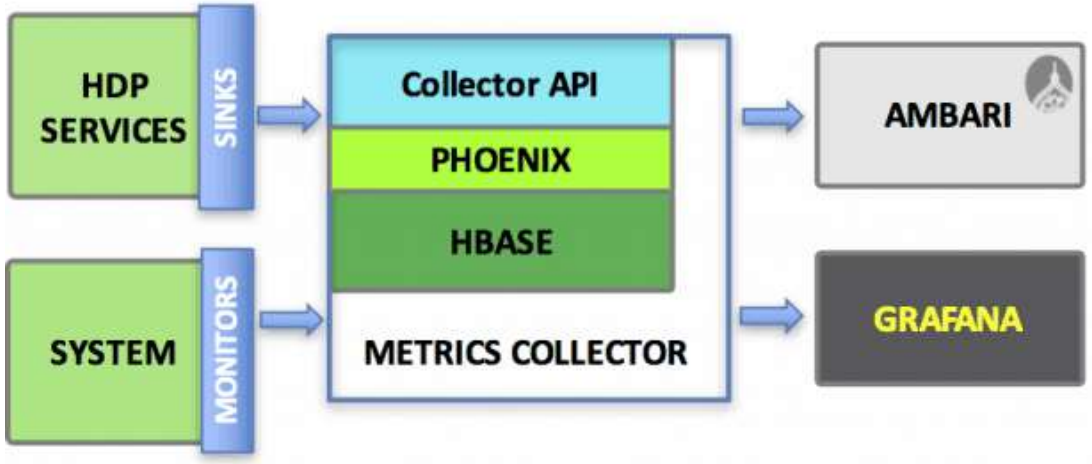

Grafana UI

METRICS LOOK GOOD

I THINK IT'S SAFE TO DEPLOY

7. Telemetry

StatsD client on each host

- Emitting metrics: CLI, Python client, Ruby client

GAUGE	INCREMENT	DECREMENT	RATE	COUNT	…

</> E.g., DataDog Client

```
from checks import AgentCheck

class HelloCheck(AgentCheck):
  def check(self, instance):
    self.gauge('hello.world', 1)
```

DataDog Examples (reference)

</> Python HTTP Check Snippet

```
start_time = time.time()
try:
  r = requests.get(url, timeout=timeout)
  end_time = time.time()
except requests.exceptions.Timeout as e:
  self.timeout_event(url, timeout, aggregation_key)
  return

if r.status_code != 200:
  self.status_code_event(url, r, aggregation_key)

timing = end_time - start_time
self.gauge('http.response_time', timing, tags=['http_check'])
```

Telemetry

- Functions: avg, min, max, rolling window, stddev
- Thresholds
- Alerts

Simple Metric

</>

```
metric = 1

try:
  hive -e "SELECT COUNT(*) FROM canary"
except SomeExceptionType:
  metric = 0

dd-report-metric -m "hive" -v $metric
```

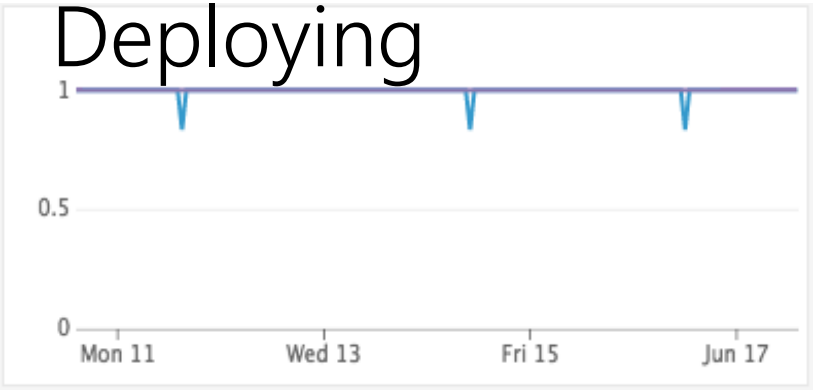

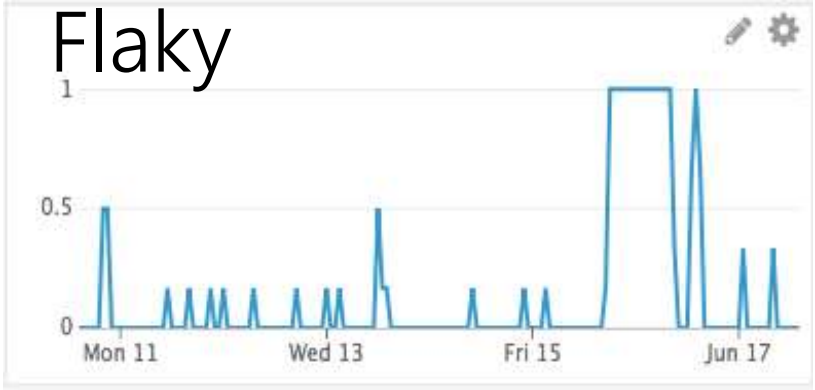

8. ETL with Apache Airflow (incubating)

Open-sourced at Airbnb by Max Beauchemin

ETL – Develop your DAGs using Python + Jinja


• Sensors: wait for data to land


• Operators galore: Bash, Hive, Presto, MySQL, etc.


• Failure tolerance, retries, notifications


• Backfill


• GUI

8. ETL with Apache Airflow (incubating)

Airflow

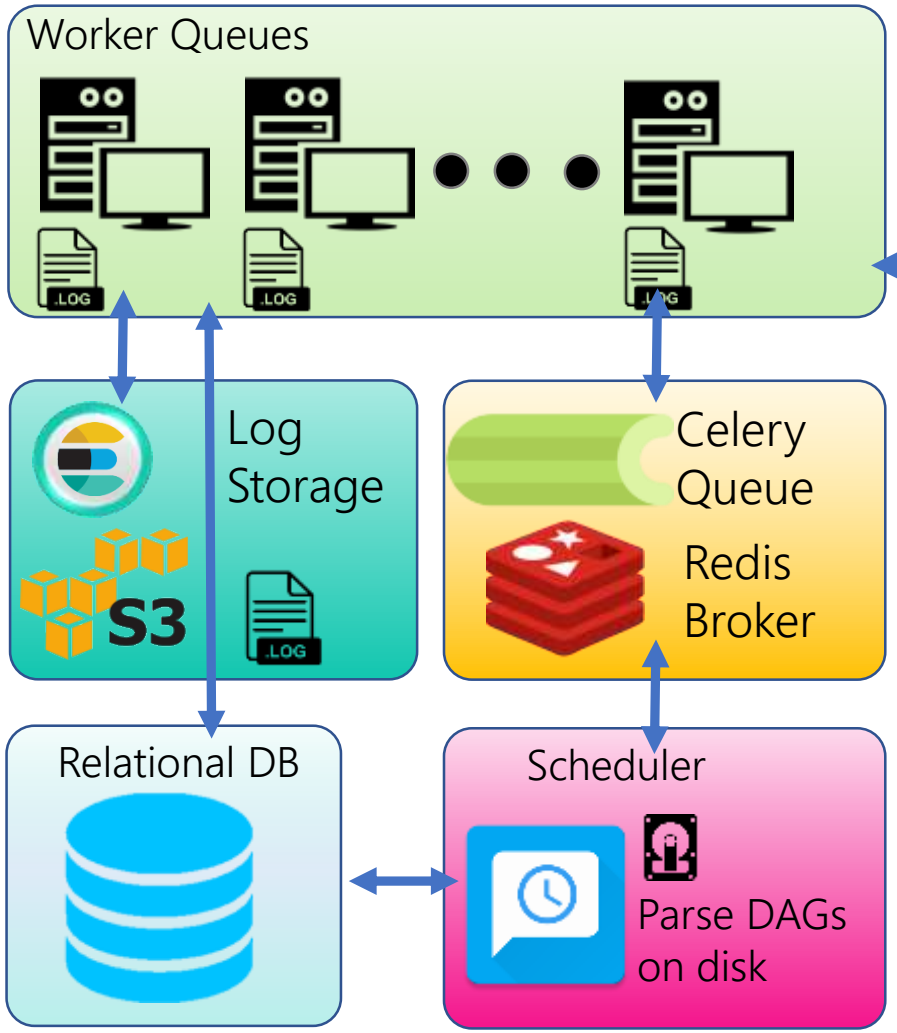

Web UI

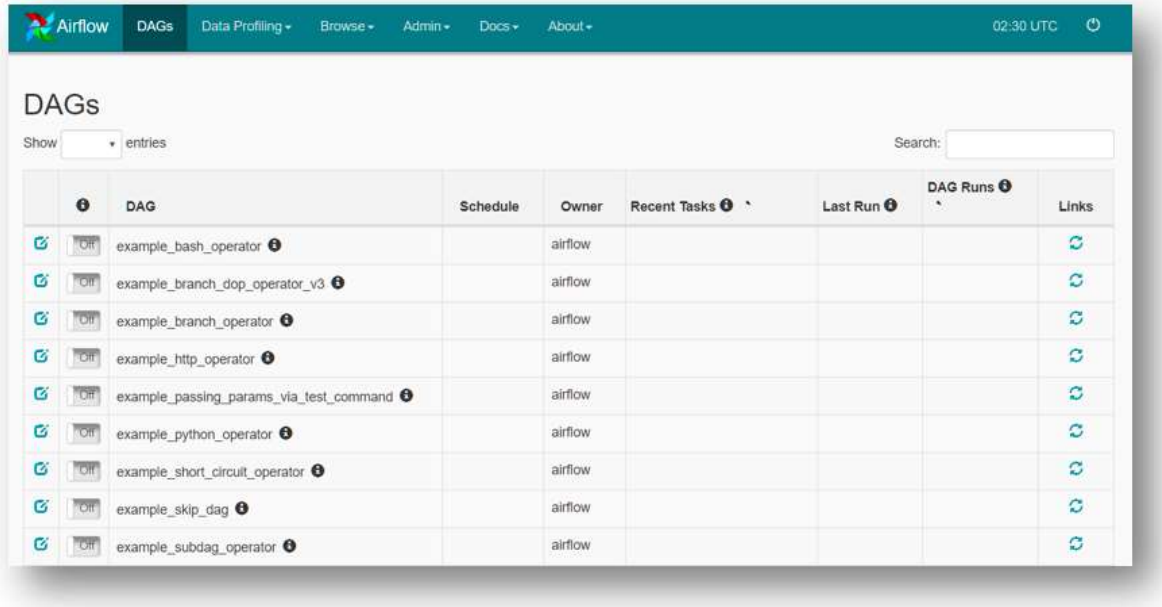

Example Bash Operator (reference)

</> DAG - Bash Operator

```
import airflow
from airflow.operators.bash_operator import BashOperator
from airflow.models import DAG


args = {
    'owner': 'airflow',
    'start_date': airflow.utils.dates.days_ago(2)
}

dag = DAG(
    dag_id='example_bash_operator',
    default_args=args,
    schedule_interval='@daily')


task = BashOperator(
    task_id='also_run_this',
    bash_command='ls -l ; echo "run_id={{ run_id }} | dag_run={{ dag_run }}"',
    dag=dag)
```

9. Dashboards with Apache Superset (incubating)

Open-sourced at Airbnb by Max Beauchemin

- Ad-hoc queries
- Connections: MySQL, SqlAlchemy, Presto, Hive, etc.
- Dashboards

9. Dashboards with Apache Superset (incubating)

Superset

Query Data

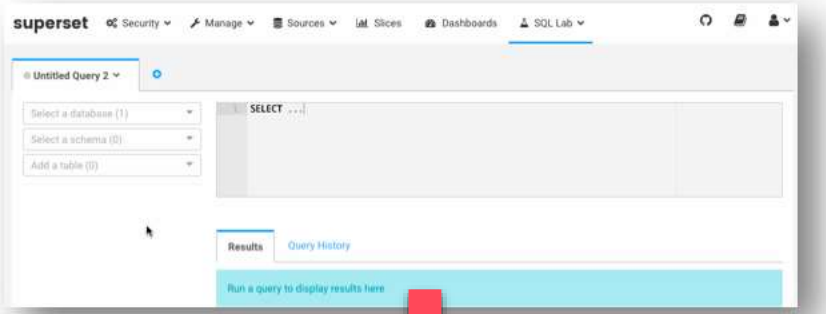

Explore Data

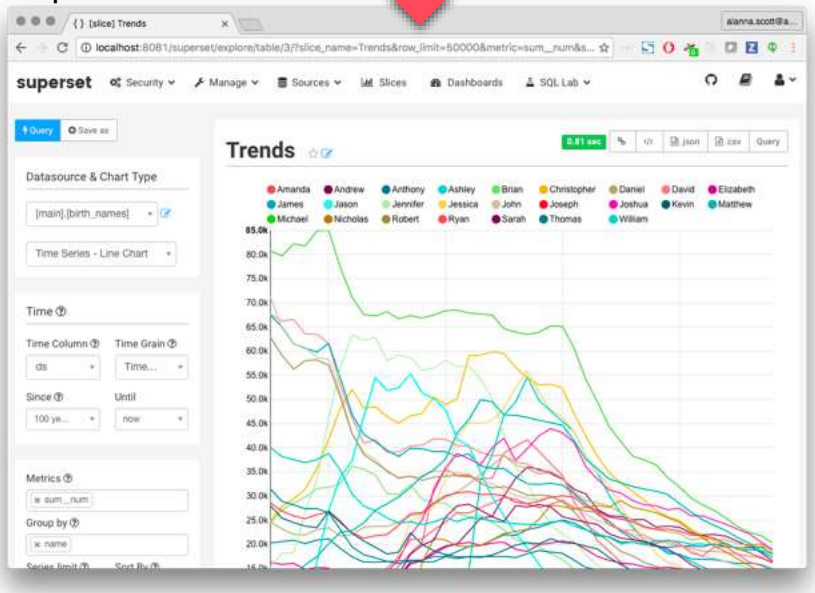

Visualize It

Share It

10. Automated Maintenance

- Rotate unhealthy nodes (too many failed volumes)
  - Decommission
  - Stop role, remove from cluster
  - Replace
- Clean space

Thank you

References

- https://medium.com/airbnb-engineering/superset-scaling-data-access-and-visual-insights-at-airbnb-3ce3e9b88a7f
- https://blog.cloudera.com/blog/2012/09/automating-your-cluster-with-cloudera-manager-api/
- https://hortonworks.com/blog/hood-ambari-metrics-grafana/
- https://community.hortonworks.com/storage/attachments/7372-ru-eu-prechecks-details.pdf
- https://community.hortonworks.com/articles/54944/stack-upgrade-pre-checks-purpose-and-remediation.html
- https://docs.datadoghq.com/developers/agent_checks/
- https://www.datadoghq.com/blog/collecting-hadoop-metrics/#namenode-and-datanode-metrics-via-jmx

Images:

- https://grafana.com/plugins/praj-ams-datasource
- https://www.estabil.is/post/o-que-e-devops
- http://warriorfitness.org/tag/top-10-list/
- https://superset.apache.org/